

```
!pip install kaggle
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.10/dist-packages (1.6.17)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.10/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: certifi>=2023.7.22 in /usr/local/lib/python3.10/dist-packages (from kaggle) (2024.8.30)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from kaggle) (4.66.5)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.10/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.0.7)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from kaggle) (6.1.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->kaggle) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.10/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.10)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.10)
```

```
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json
```

```
cp: cannot stat 'kaggle.json': No such file or directory
chmod: cannot access '/root/.kaggle/kaggle.json': No such file or directory
```

```
# ApI to fetch the dataset from Kaggle
!kaggle datasets download -d milobele/sentiment140-dataset-1600000-tweets
```

```
Dataset URL: https://www.kaggle.com/datasets/milobele/sentiment140-dataset-1600000-tweets
License(s): unknown
Downloading sentiment140-dataset-1600000-tweets.zip to /content
 96% 78.0M/81.0M [00:01<00:00, 89.4MB/s]
100% 81.0M/81.0M [00:01<00:00, 83.3MB/s]
```

```
# extracting the compressed dataset
from zipfile import ZipFile
import os
dataset = '/content/sentiment140-dataset-1600000-tweets.zip'

# Check if the dataset file exists
if os.path.exists(dataset):
    with ZipFile(dataset, 'r') as zip:
        zip.extractall() # Extract all contents to the current directory
        print('The dataset is extracted')
else:
    print(f'The file {dataset} does not exist.')
```

```
The dataset is extracted
```

## Importing Libraries

```
# utilites
import pandas as pd
import numpy as np
import re
# plotting
import seaborn as sns
from wordcloud import WordCloud
import matplotlib.pyplot as plt
# nltk
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from nltk.stem import PorterStemmer
from nltk.stem import WordNetLemmatizer
# sklearn
from sklearn.svm import LinearSVC
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
```

```

from sklearn.naive_bayes import BernoulliNB
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB

```

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.

```

```

#printing the stopwaords in English
stopwordlist=stopwords.words('english')
print(len(stopwordlist))
print(stopwordlist)

```

```

179
['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'your', 'y

```

## Data Processing

```

#load the data from csv file to pandas dataframe
df=pd.read_csv('/content/training.1600000.processed.noemoticon.csv',encoding='latin-1')

```

```
df.shape
```

```
(1599999, 6)
```

```
df.head()
```

```

0  1467810369  Mon Apr 06 22:19:45 PDT 2009  NO_QUERY  _TheSpecialOne_  @switchfoot http://twitpic.com/2y1zl - Awww, that's a bummer. You shoulda got David Carr of Third Day to do it. ;D

0  0  1467810672  Mon Apr 06 22:19:49 PDT 2009  NO_QUERY  scotthamilton  is upset that he can't update his Facebook by ...

1  0  1467810917  Mon Apr 06 22:19:53 PDT 2009  NO_QUERY  mattycus  @Kenichan I dived many times for the ball. Man...

2  0  1467811184  Mon Apr 06 22:19:57 PDT 2009  NO_QUERY  ElleCTF  my whole body feels itchy and like its on fire

3  0  1467811193  Mon Apr 06 22:19:57 PDT 2009  NO_QUERY  Karoli  @nationwiderclass no, it's not behaving at all

```

```

#naming the columns and reading again
column_names = ['target','ids','date','flag','user','text']
df = pd.read_csv('/content/training.1600000.processed.noemoticon.csv',header=None,encoding='latin-1',names=column_names)

```

```
df.head()
```

```

target      ids      date      flag      user      text

0      0  1467810369  Mon Apr 06 22:19:45 PDT 2009  NO_QUERY  _TheSpecialOne_  @switchfoot http://twitpic.com/2y1zl - Awww, t...

1      0  1467810672  Mon Apr 06 22:19:49 PDT 2009  NO_QUERY  scotthamilton  is upset that he can't update his Facebook by ...

2      0  1467810917  Mon Apr 06 22:19:53 PDT 2009  NO_QUERY  mattycus  @Kenichan I dived many times for the ball. Man...

```

```

# Columns/ feature in data
df.columns

```

```
Index(['target', 'ids', 'date', 'flag', 'user', 'text'], dtype='object')
```

```
# Length of the dataset
print('length of data is',len(df))
```

```
➦ length of data is 1600000
```

```
# Shape of data
df.shape
```

```
➦ (1600000, 6)
```

```
# Data information
df.info()
```

```
➦ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1600000 entries, 0 to 1599999
Data columns (total 6 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   target  1600000 non-null  int64  
 1   ids     1600000 non-null  int64  
 2   date    1600000 non-null  object  
 3   flag    1600000 non-null  object  
 4   user    1600000 non-null  object  
 5   text    1600000 non-null  object  
dtypes: int64(2), object(4)
memory usage: 73.2+ MB
```

```
# checking for null values
df.isnull().sum()
```

```
➦
```

	0
target	0
ids	0
date	0
flag	0
user	0
text	0

```
np.sum(df.isnull().any(axis=1))
```

```
➦ 0
```

```
# Check unique target values
df['target'].unique()
```

```
➦ array([0, 4])
```

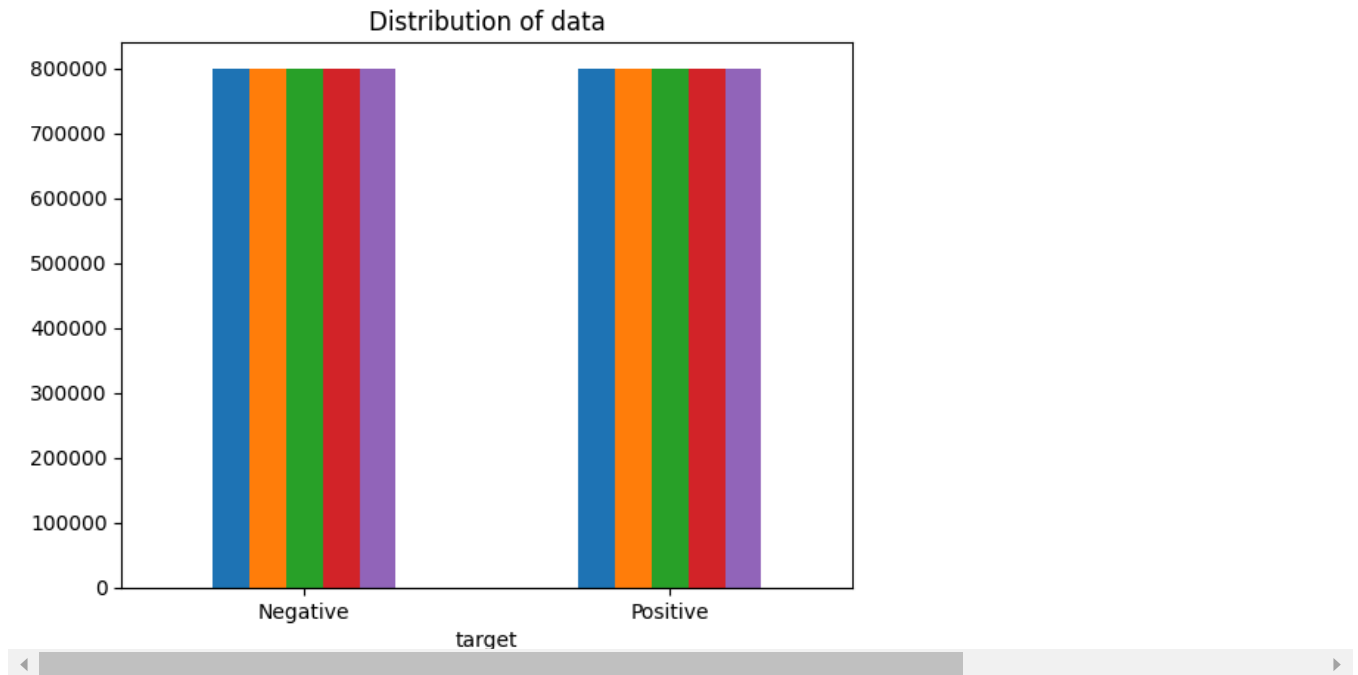
```
# Checking the number of target values
df['target'].nunique()
```

```
➦ 2
```

## Data Visualization of Target Variables

```
# plotting the distribution for dataset.
ax= df.groupby('target').count().plot(kind='bar', title='Distribution of data',legend=False)
ax.set_xticklabels(['Negative','Positive'], rotation=0)
```

```
🔗 [Text(0, 0, 'Negative'), Text(1, 0, 'Positive')]
```



The target variables of contains positive and negative vales are equal

### Data Preprocessing

In the above-given problem statement, before training the model, we performed various pre-processing steps on the dataset that mainly dealt with removing stopwords, removing special characters like emojis, hashtags, etc. The text document is then converted into lowercase for better generalization.

Subsequently, the punctuations were cleaned and removed, thereby reducing the unnecessary noise from the dataset. After that, we also removed the repeating characters from the words along with removing the URLs as they do not have any significant importance. At last, we then performed Stemming(reducing the words to their derived stems) and Lemmatization(reducing the derived words to their root form, known as lemma) for better results.

Selecting the Text and Target column for our further analysis

```
data= df[['text','target']]
```

Replacing the values to easy understanding (Assigning 1 to Positive sentiment 4)

```
data['target']=data['target'].replace(4,1)
```

```
🔗 <ipython-input-22-784d662c4464>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-data['target']=data['target'].replace(4,1)  
data['target']=data['target'].replace(4,1)

```
# Printing unique values of target variables  
data['target'].unique()
```

```
🔗 array([0, 1])
```

```
# Seprating positive and negative tweets  
data_pos=data[data['target']==1]  
data_neg=data[data['target']==0]
```

```
# Taking one-fourth of the data so we can run it on our machine easily
data_pos=data_pos.iloc[:len(data_pos)]
data_neg=data_neg.iloc[:len(data_neg)]
```

```
print(data_pos.shape)
print(data_neg.shape)
```

```
(800000, 2)
(800000, 2)
```

```
# Combining positive and negative tweets
dataset= pd.concat([data_pos,data_neg])
```

```
# Convert all text to lowercase
dataset['text']= dataset['text'].str.lower()
print('This is head of text data\n',dataset['text'].head())
print('\n')
```

```
print('This is tail of text data\n',dataset['text'].tail())
```

```
This is head of text data
800000      i love @health4uandpets u guys r the best!!
800001  im meeting up with one of my besties tonight! ...
800002  @darealsunisakim thanks for the twitter add, s...
800003  being sick can be really cheap when it hurts t...
800004  @lovesbrooklyn2 he has that effect on everyone
Name: text, dtype: object
```

```
This is tail of text data
799995  sick spending my day laying in bed listening ...
799996  gmail is down?
799997  rest in peace farrah! so sad
799998  @eric_urbane sounds like a rival is flagging y...
799999  has to resit exams over summer... wishes he w...
Name: text, dtype: object
```

## Removing the stopwordslist from the tweets text columns

```
stopwordlist = ['a', 'about', 'above', 'after', 'again', 'ain', 'all', 'am', 'an',
'and', 'any', 'are', 'as', 'at', 'be', 'because', 'been', 'before',
'being', 'below', 'between', 'both', 'by', 'can', 'd', 'did', 'do',
'does', 'doing', 'down', 'during', 'each', 'few', 'for', 'from',
'further', 'had', 'has', 'have', 'having', 'he', 'her', 'here',
'hers', 'herself', 'him', 'himself', 'his', 'how', 'i', 'if', 'in',
'into', 'is', 'it', 'its', 'itself', 'just', 'll', 'm', 'ma',
'me', 'more', 'most', 'my', 'myself', 'now', 'o', 'of', 'on', 'once',
'only', 'or', 'other', 'our', 'ours', 'ourselves', 'out', 'own', 're', 's', 'same', 'she', "shes", 'should', "s
't', 'than', 'that', "thatll", 'the', 'their', 'theirs', 'them',
'themselves', 'then', 'there', 'these', 'they', 'this', 'those',
'through', 'to', 'too', 'under', 'until', 'up', 've', 'very', 'was',
'we', 'were', 'what', 'when', 'where', 'which', 'while', 'who', 'whom',
'why', 'will', 'with', 'won', 'y', 'you', 'youd', "youll", "youre",
"youve", 'your', 'yours', 'yourself', 'yourselves', 'u', 'r']
```

```
STOPWORDS = set(stopwordlist) # set of stop words
def clean_stopwordslist(text):
    return " ".join([word for word in str(text).split() if word not in STOPWORDS])
dataset['text'] = dataset['text'].apply(lambda text: clean_stopwordslist(text))
print('This is head of text data\n',dataset['text'].head())
print('\n')
```

```
print('This is tail of text data\n',dataset['text'].tail())
```

```
This is head of text data
800000      love @health4uandpets guys best!!
800001  im meeting one besties tonight! cant wait!! - ...
800002  @darealsunisakim thanks twitter add, sunisa! g...
800003  sick really cheap hurts much eat real food plu...
800004  @lovesbrooklyn2 effect everyone
Name: text, dtype: object
```

```

This is tail of text data
799995    sick spending day laying bed listening @taylor...
799996                                           gmail down?
799997                                           rest peace farrah! sad
799998    @eric_urbane sounds like rival flagging ads. n...
799999    resit exams over summer... wishes worked harde...
Name: text, dtype: object

```

```

STOPWORDS = set(stopwords.words('english')) # set of stop words
def clean_stopwords(text):
    return " ".join([word for word in str(text).split() if word not in STOPWORDS])
dataset['text'] = dataset['text'].apply(lambda text: clean_stopwords(text))
print('This is head of text data\n',dataset['text'].head())
print('\n')

print('This is tail of text data\n',dataset['text'].tail())

```

```

➡ This is head of text data
800000    love @health4uandpets guys best!!
800001    im meeting one besties tonight! cant wait!! - ...
800002    @darealsunisakim thanks twitter add, sunisa! g...
800003    sick really cheap hurts much eat real food plu...
800004    @lovesbrooklyn2 effect everyone
Name: text, dtype: object

```

```

This is tail of text data
799995    sick spending day laying bed listening @taylor...
799996                                           gmail down?
799997                                           rest peace farrah! sad
799998    @eric_urbane sounds like rival flagging ads. m...
799999    resit exams summer... wishes worked harder fir...
Name: text, dtype: object

```

```
dataset['text'][800001]
```

```
➡ 'im meeting one besties tonight! cant wait!! - ...'
```

```

def cleaning_urls(data):
    return re.sub('((www.[^s]+)|(https?://[^\s]+))',' ',data)
dataset['text'] = dataset['text'].apply(lambda x: cleaning_urls(x))
dataset['text'].head()

```

```

➡

```

	text
800000	love @health4uandpets guys best!!
800001	im meeting one besties tonight! cant wait!! - ...
800002	@darealsunisakim thanks twitter add, sunisa! g...
800003	sick really cheap hurts much eat real food plu...
800004	@lovesbrooklyn2 effect everyone

```

Name: text, dtype: object

```

## Removing punctuations

```

import string
english_punctuations = string.punctuation
punctuations_list = english_punctuations

def cleaning_punctuations(text):
    translator = str.maketrans('', '', punctuations_list)
    return text.translate(translator)

dataset['text']= dataset['text'].apply(lambda x: cleaning_punctuations(x))
print('This is head of text data\n',dataset['text'].head())
print('\n')

print('This is tail of text data\n',dataset['text'].tail())

```

```

➡ This is head of text data
800000          love health4uandpets guys best
800001  im meeting one besties tonight cant wait  girl...
800002  darealsunisakim thanks twitter add sunisa got ...
800003  sick really cheap hurts much eat real food plu...
800004          lovesbrooklyn2 effect everyone
Name: text, dtype: object

```

```

This is tail of text data
799995  sick spending day laying bed listening taylors...
799996                                     gmail down
799997                                     rest peace farrah sad
799998  ericurbane sounds like rival flagging ads much...
799999  resit exams summer wishes worked harder first ...
Name: text, dtype: object

```

## Removing repeating characters

```

# def cleaning_repeating_char(text):
#     return re.sub(r'(.+)+', r'1', text)
# dataset['text'] = dataset['text'].apply(lambda x: cleaning_repeating_char(x))

# print('This is head of text data\n',dataset['text'].head())
# print('\n')

# print('This is tail of text data\n',dataset['text'].tail())

```

## Removing numerical values from the text columns

```

def cleaning_numerical(data):
    return re.sub('[0-9]+','', data)
dataset['text'] = dataset['text'].apply(lambda x: cleaning_numerical(x))
print('This is head of text data\n',dataset['text'].head())
print('\n')

print('This is tail of text data\n',dataset['text'].tail())

```

```

➡ This is head of text data
800000          love healthuandpets guys best
800001  im meeting one besties tonight cant wait  girl...
800002  darealsunisakim thanks twitter add sunisa got ...
800003  sick really cheap hurts much eat real food plu...
800004          lovesbrooklyn effect everyone
Name: text, dtype: object

```

```

This is tail of text data
799995  sick spending day laying bed listening taylors...
799996                                     gmail down
799997                                     rest peace farrah sad
799998  ericurbane sounds like rival flagging ads much...
799999  resit exams summer wishes worked harder first ...
Name: text, dtype: object

```

## Getting tokenization of tweets test

```

# from nltk.tokenize import RegexpTokenizer
# tokenizer = RegexpTokenizer(r'\w+')
# dataset['text'] = dataset['text'].apply(tokenizer.tokenize)
# print('This is head of text data\n',dataset['text'].head())
# print('\n')

# print('This is tail of text data\n',dataset['text'].tail())
# dataset.head()

```

## ✓ Stemming

- Reducing the words to their derived stems

- stemming is the process of reducing a word to its word stem that affixes to suffixes and prefixes or to the roots of words known as a lemma.

eg.

- Porter Stemmer:

Original: running, Stemmed: run

Original: runner, Stemmed: runner

Original: ran, Stemmed: ran

```
from nltk.stem import PorterStemmer

port_stemmer = PorterStemmer()

def stemming_on_text(data):

    # Call the stem method of port_stemmer object
    text=[port_stemmer.stem(word) for word in data]
    return data

dataset['text'] = dataset['text'].apply(lambda x: stemming_on_text(x))
print('This is head of text data\n',dataset['text'].head())
print('\n')

print('This is tail of text data\n',dataset['text'].tail())
```

```
↗ This is head of text data
800000          love healthuandpets guys best
800001  im meeting one besties tonight cant wait girl...
800002  darealsunisakim thanks twitter add sunisa got ...
800003  sick really cheap hurts much eat real food plu...
800004          lovesbrooklyn effect everyone
Name: text, dtype: object
```

```
This is tail of text data
799995  sick spending day laying bed listening taylors...
799996          gmail down
799997          rest peace farrah sad
799998  ericurbane sounds like rival flagging ads much...
799999  resit exams summer wishes worked harder first ...
Name: text, dtype: object
```

dataset.head()

```
↗
```

	text	target
800000	love healthuandpets guys best	1
800001	im meeting one besties tonight cant wait girl...	1
800002	darealsunisakim thanks twitter add sunisa got ...	1
800003	sick really cheap hurts much eat real food plu...	1
800004	lovesbrooklyn effect everyone	1

## ▼ Lemmatizer

lemmatization is a process that transforms words into their base or root form.

For example, the words "running," "ran," and "runs" would all be lemmatized to the base form "run."

```
import nltk
!pip install nltk
nltk.download('wordnet') # downloads the 'wordnet' resource

lm= nltk.WordNetLemmatizer()
```



```
def lemmatizer_on_text(data):

    text=[lm.lemmatize(word) for word in data]
    return data


dataset['text'] = dataset['text'].apply(lambda x: lemmatizer_on_text(x))
dataset['text'].head()
```

 [Show hidden output](#)

Seprating input features and label

```
X=data.text
y=data.target
```

```
print(len(X))
print(len(y))
```

 1600000  
1600000

Plot a cloud of words for Negative tweets

```
!pip install wordcloud
import matplotlib.pyplot as plt
from wordcloud import WordCloud

data_neg = data['text']
plt.figure(figsize = (20,20))
wc = WordCloud(max_words = 1000 , width = 1600 , height = 800,
               collocations=False).generate(" ".join(data_neg))
plt.imshow(wc)
```





```

(0, 471581)    0.4369096022969522
(0, 234798)    0.16395910802365576
(0, 147252)    0.36841582248659926
(1, 253367)    0.27553894838443566
(1, 42112)     0.2728120613981423
(1, 235103)    0.09464334482888374
:              :
(1119998, 222693) 0.2304746111612065
(1119998, 509124) 0.24669180387566617
(1119998, 87499)  0.18946206616692834
(1119998, 464148) 0.23129988464226583
(1119998, 176365) 0.21214435066754267
(1119998, 8855)   0.23134648116874723
(1119998, 326113) 0.16183601296231373
(1119998, 465035) 0.16960775464145403
(1119998, 294809) 0.11739269358984537
(1119998, 362261) 0.08948076217276431
(1119998, 499385) 0.21069286874004575
(1119998, 72701)  0.09971911099934253
(1119998, 482837) 0.06074447605076158
(1119999, 519706) 0.5278675821833199
(1119999, 456245) 0.3478118774961449
(1119999, 174609) 0.3433527141178553
(1119999, 432536) 0.3104178885358356
(1119999, 311883) 0.38437633361830725
(1119999, 218638) 0.22797175718256743
(1119999, 228370) 0.21843869069341512
(1119999, 199142) 0.2053826704680457
(1119999, 477904) 0.1872905467393014
(1119999, 442759) 0.16530147597641823
(1119999, 183969) 0.14905270490387504
(1119999, 482837) 0.10908499072351457

```

```
print(X_test)
```



```

(0, 531750)    0.2711284684055641
(0, 360763)    0.3569742620291956
(0, 219592)    0.3094720116021608
(0, 112650)    0.6935524442506389
(0, 41366)     0.47145605413192115
(1, 491934)    0.26272642629141696
(1, 397927)    0.5537642962993318
(1, 353350)    0.1821674564234399
(1, 340109)    0.12165514926170792
(1, 327117)    0.5777500415031038
(1, 210526)    0.15352844250975373
(1, 183969)    0.13678612518709593
(1, 122912)    0.25412731858906384
(1, 57528)     0.3683445945541971
(2, 517989)    0.1169550299777069
(2, 514018)    0.24573130779568927
(2, 480589)    0.1388282281703462
(2, 479455)    0.45423643136881686
(2, 439606)    0.2989217947550147
(2, 365704)    0.12764144358248303
(2, 364183)    0.15802787599390913
(2, 348686)    0.17452087057365034
(2, 289219)    0.12936177061497478
(2, 280839)    0.24355922926085932
(2, 274036)    0.1422062689008427
:              :
(479998, 433698) 0.1946345013120972
(479998, 406340) 0.1874760039927674
(479998, 398948) 0.16729738027613278
(479998, 359104) 0.11973399928849832
(479998, 355643) 0.14252978494477592
(479998, 340109) 0.10108550767052475
(479998, 311767) 0.33447148698856843
(479998, 234240) 0.11022383159285648
(479998, 229713) 0.11394099616401644
(479998, 178492) 0.3873099686878723
(479998, 166696) 0.25937483035549874
(479998, 142300) 0.30277535266317546
(479998, 137779) 0.2017458897452608
(479998, 91865)  0.2558948512702573
(479998, 76739)  0.35610291871454
(479998, 74996)  0.20796953282549438
(479998, 72701)  0.1365517256581464
(479998, 46177)  0.1474823819297551
(479999, 509128) 0.3236754463153666

```

(479999, 484926)	0.22300389028219178
(479999, 411042)	0.48410197011441675
(479999, 304933)	0.5939933126948371
(479999, 235103)	0.1484633379145862
(479999, 188192)	0.45814054184969527
(479999, 183969)	0.16243946129446946

## ✓ Function for Model Evaluation

After training the model , we apply the evaluation measures to check how the model is performing. Accordingly , we use the following evaluation parameters to check the performance of the models respectively

- Accuracy Score
- Confusion Matrix with Plot
- ROC- AUC Curve

```
import zipfile # import zipfile module
from sklearn.naive_bayes import BernoulliNB
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import builtins # import the builtins module

def model_Evaluate(model):
    # Predict values for Test dataset
    y_pred = model.predict(X_test)

    # Print the evaluation metrics for the dataset.
    print(classification_report(y_test, y_pred))

    # Print the accuracy score
    accuracy_test = accuracy_score(y_test, y_pred) # Calculate accuracy
    print(f'Accuracy: {accuracy_test:.2f}') # Print accuracy score

    # Compute and plot the Confusion matrix
    cf_matrix = confusion_matrix(y_test, y_pred)
    categories = ['Negative', 'Positive']
    group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
    group_percentages = ['{0:.2%}'.format(value) for value in cf_matrix.flatten() / np.sum(cf_matrix)]

    # use the built-in zip function. The user probably overwrote the zip function with a ZipFile object
    labels = [f'{v1}\n{v2}' for v1, v2 in builtins.zip(group_names, group_percentages)]
    labels = np.asarray(labels).reshape(2,2)

    sns.heatmap(cf_matrix, annot = labels, cmap = 'Blues', fmt = '',
                xticklabels = categories, yticklabels = categories)
    plt.xlabel("Predicted values", fontdict = {'size':14}, labelpad = 10)
    plt.ylabel("Actual values", fontdict = {'size':14}, labelpad = 10)
    plt.title ("Confusion Matrix", fontdict = {'size':18}, pad = 20)
```

## Model Building

In the problem statement, we have used three different models respectively:

- Bernoulli Naive Bayes Classifier
- SVM (Support Vector Machine)
- Logistic Regression

The idea behind choosing these models is that we want to try all the classifiers on the dataset ranging from simple ones to complex models, and then try to find out the one which gives the best performance among them


## ✓ Model -1 Bernoulli Naive Bayes Classifier

```
BNBmodel = BernoulliNB()
BNBmodel.fit(X_train, y_train)
model_Evaluate(BNBmodel)
y_pred1 = BNBmodel.predict(X_test)
```

	precision	recall	f1-score	support
0	0.77	0.81	0.79	239361
1	0.80	0.76	0.78	240639
accuracy			0.78	480000
macro avg	0.78	0.78	0.78	480000
weighted avg	0.78	0.78	0.78	480000

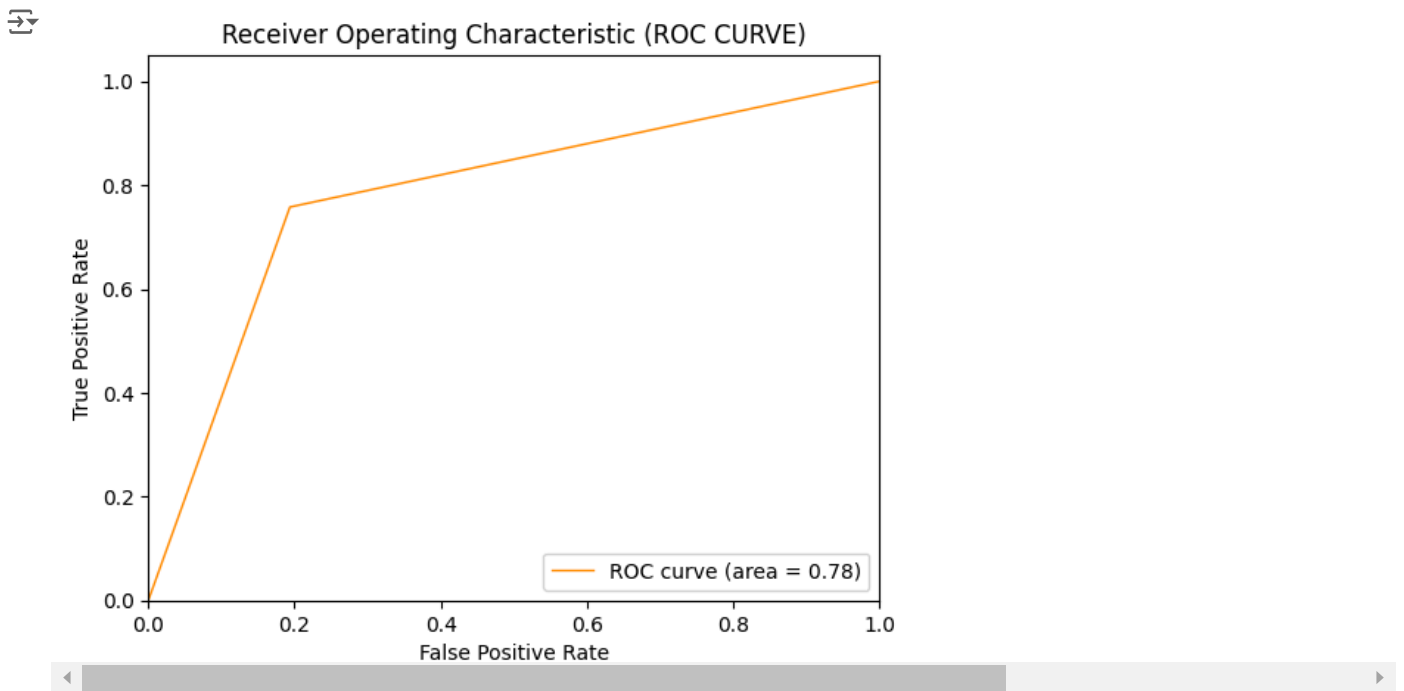
Accuracy: 0.78

Confusion Matrix



## ✓ Plot the ROC-AUC Curve for model-1

```
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred1)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=1, label='ROC curve (area = %0.2f)' % roc_auc)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC CURVE)')
plt.legend(loc="lower right")
plt.show()
```



## Model-2 Support Vector Machine


```
SVCmodel = LinearSVC()  
SVCmodel.fit(X_train, y_train)  
model_Evaluate(SVCmodel)  
y_pred2 = SVCmodel.predict(X_test)
```

/usr/local/lib/python3.10/dist-packages/sklearn/svm/\_classes.py:32: FutureWarning: The default value of `dual` will change from `True` to `False` in 1.5. To suppress this warning, please use `dual=False` when creating the `SVC` object.

	precision	recall	f1-score	support
0	0.80	0.79	0.79	239361
1	0.79	0.80	0.80	240639
accuracy			0.80	480000
macro avg	0.80	0.80	0.80	480000
weighted avg	0.80	0.80	0.80	480000


Accuracy: 0.80

## Confusion Matrix



## ✓ Plot the ROC-AUC Curve for model-2

```
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred2)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=1, label='ROC curve (area = %0.2f)' % roc_auc)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC CURVE)')
plt.legend(loc="lower right")
plt.show()
```



## ✓ Model-3 Logistic Regression

```
LRmodel = LogisticRegression(C = 2,max_iter= 1000, n_jobs=-1)
LRmodel.fit(X_train, y_train)
model_Evaluate(LRmodel)
y_pred3 = LRmodel.predict(X_test)
```







	precision	recall	f1-score	support
0	0.81	0.79	0.80	239361
1	0.80	0.81	0.80	240639
accuracy			0.80	480000
macro avg	0.80	0.80	0.80	480000
weighted avg	0.80	0.80	0.80	480000

Accuracy: 0.80

## Confusion Matrix



Plot the ROC-AUC Curve for model-3



```
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred3)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=1, label='ROC curve (area = %0.2f)' % roc_auc)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC CURVE)')
plt.legend(loc="lower right")
plt.show()
```

